

---

# Programming Map/Reduce

Hadoop Docs:

<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Tools:

- Eclipse (w/plugin <https://github.com/winghc/hadoop2x-eclipse-plugin> )
- IntelliJ Idea

Hadoop on Docker ([https://www.cloudera.com/documentation/enterprise/5-6-x/topics/quickstart\\_docker\\_container.html](https://www.cloudera.com/documentation/enterprise/5-6-x/topics/quickstart_docker_container.html) ):

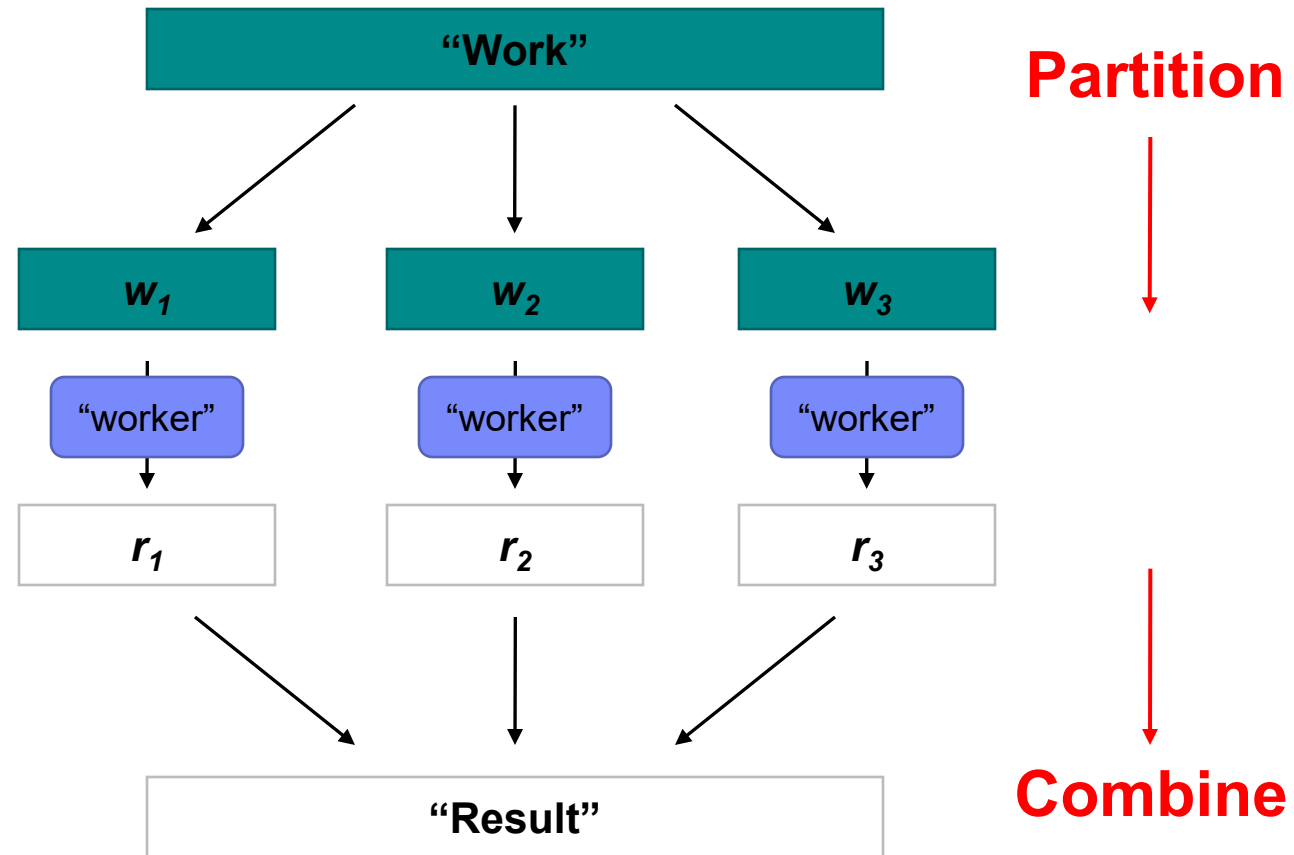
```
run --name hadoop -p 8888:8888 -p 8000:80 --privileged=true --hostname=quickstart.cloudera -t -i -d -v D:/:/shared/d -p 8088:8088 cloudera/quickstart:latest /usr/bin/docker-quickstart
```

Demos:

Python: <http://www.michael-noll.com/tutorials/writing-an-hadoop-mapreduce-program-in-python/>

Java: WordCount (in class)

## Divide and Conquer

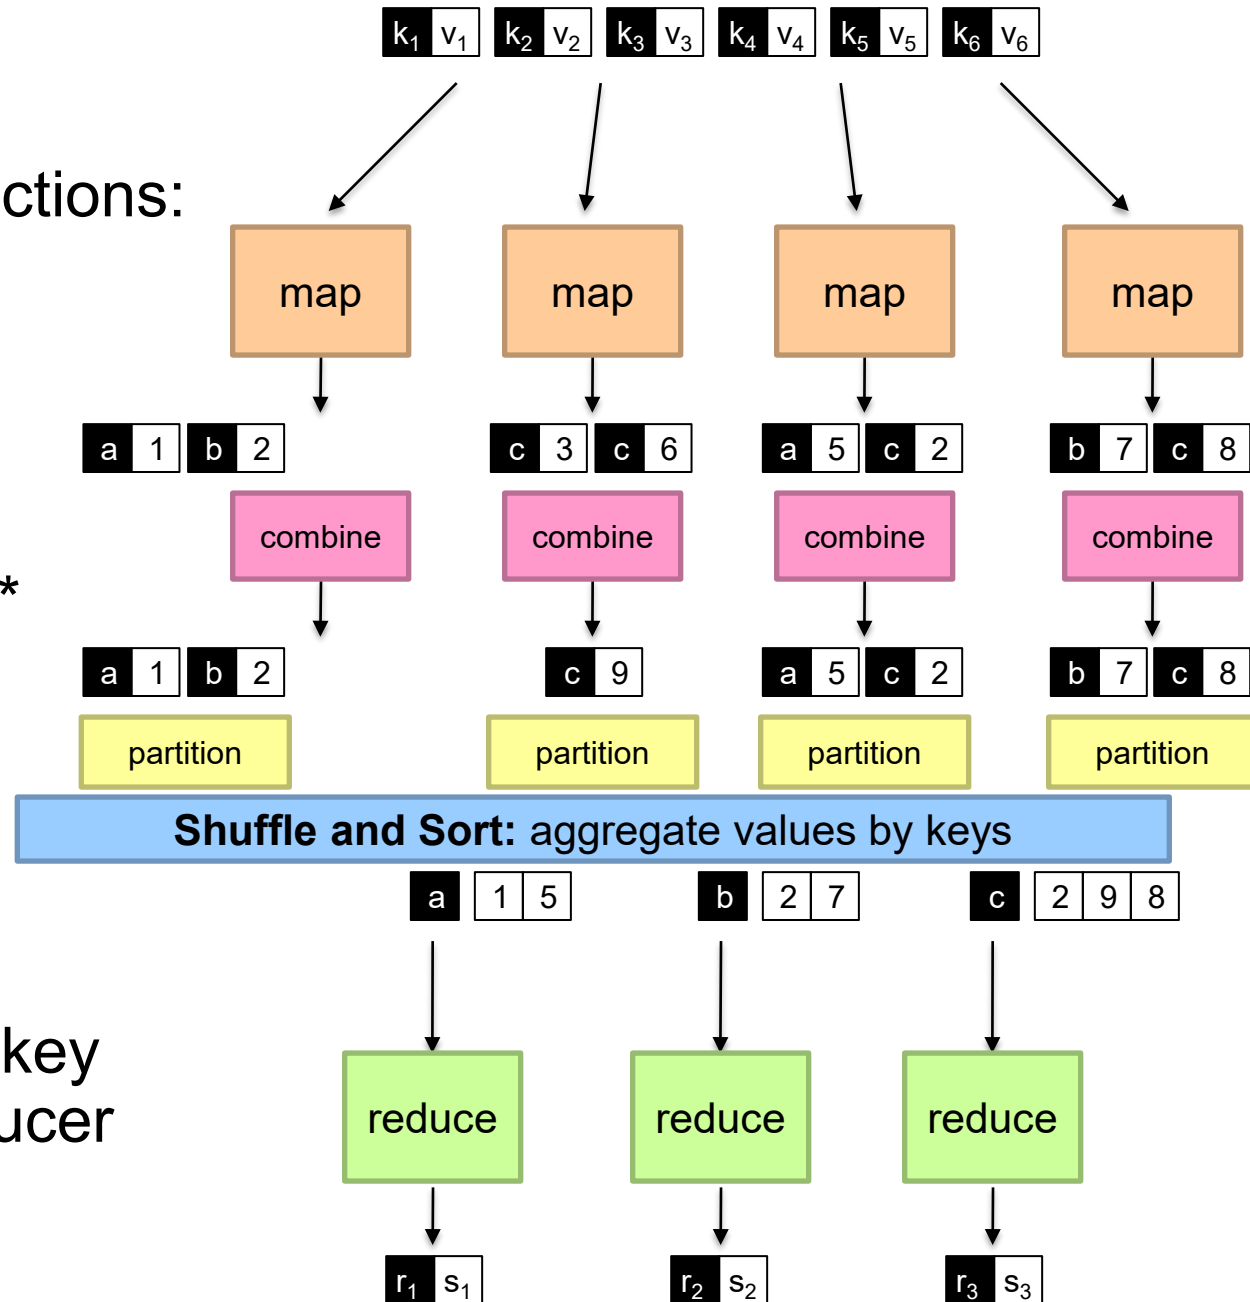


# MapReduce:

- Programmers specify functions:

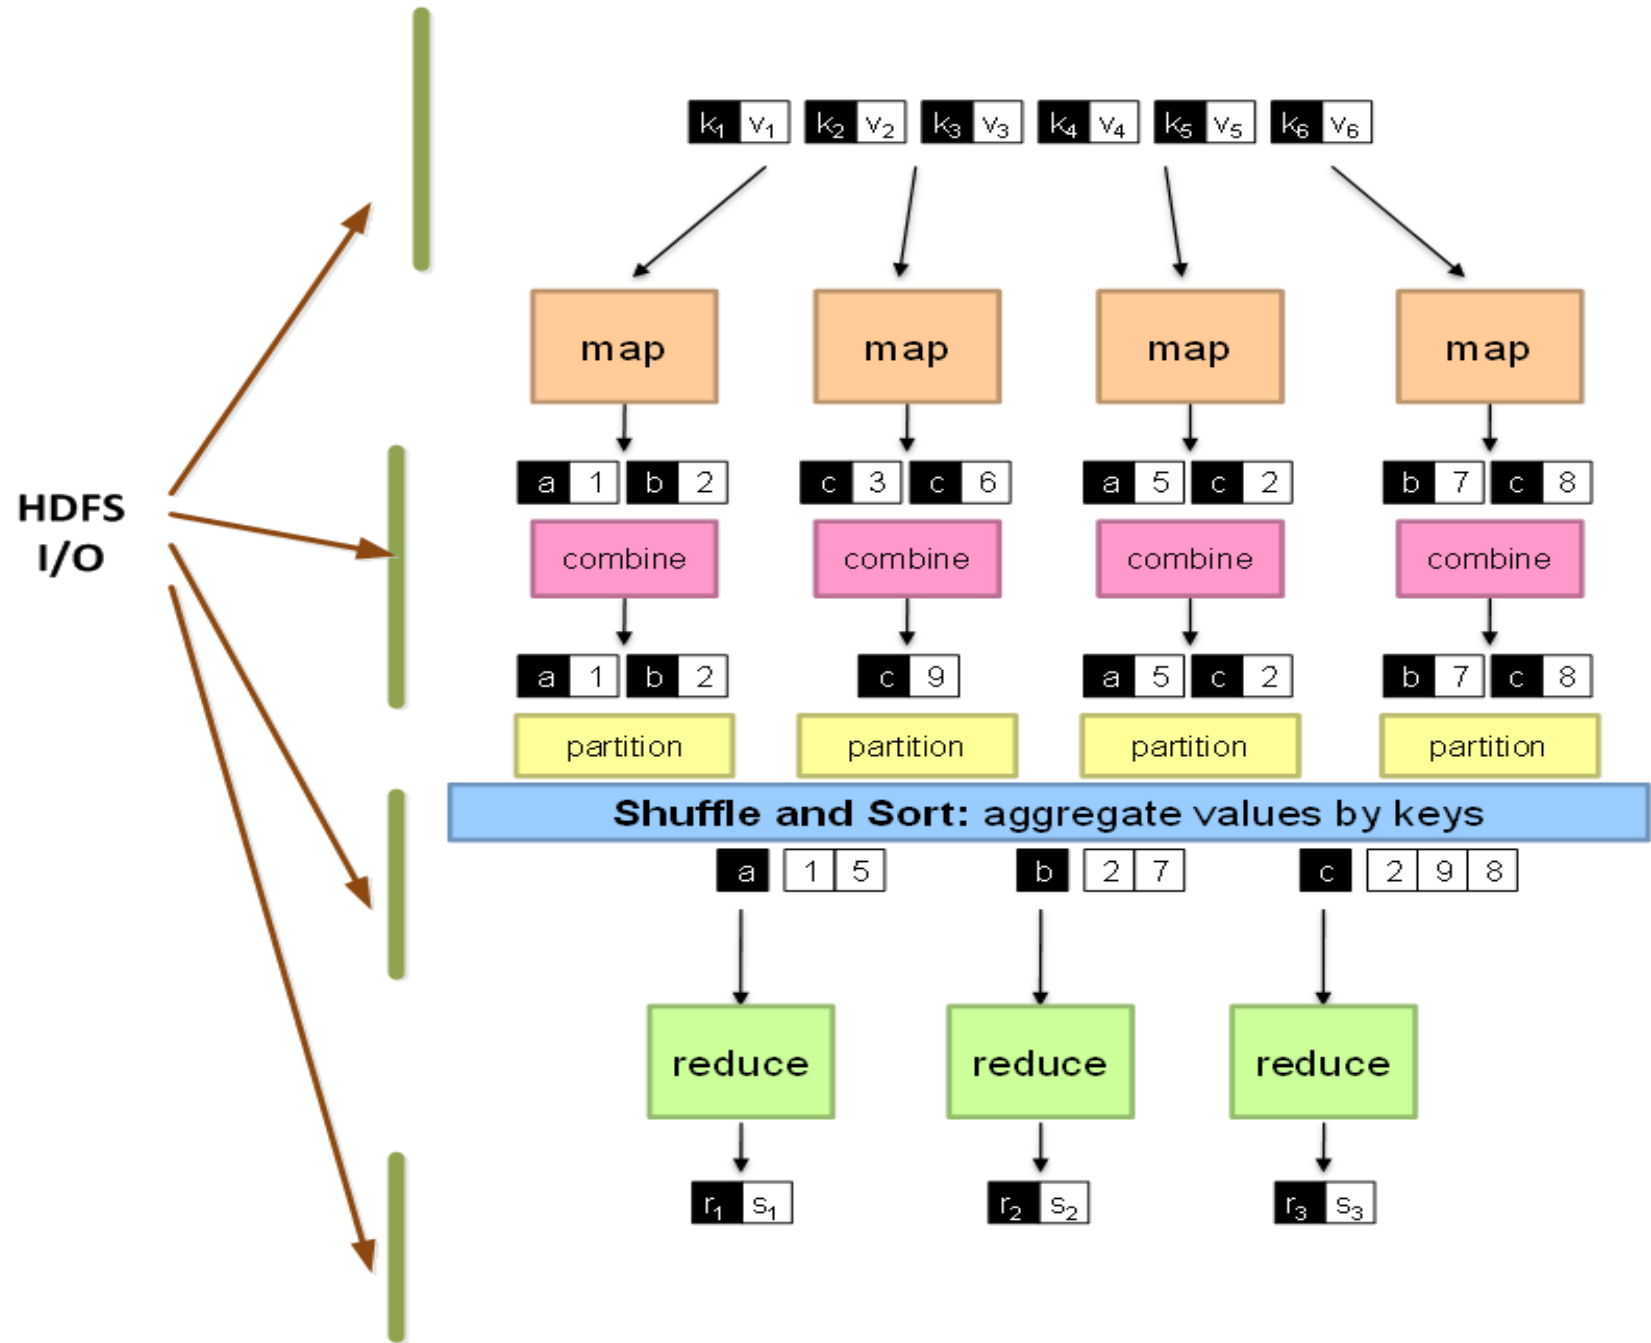
**map**  $(k, v) \rightarrow \langle k', v' \rangle^*$

**reduce**  $(k', v') \rightarrow \langle k', v' \rangle^*$



All values with the same key  
are sent to the same reducer

# MapReduce:



# MapReduce

## Divided in two phases

- Map phase
- Reduce phase
- **Both phases use key-value pairs as input and output**
- **The implementer provides map and reduce functions**
- **MapReduce framework orchestrates splitting, and distributing of Map and Reduce phases**

Most of the pieces can be easily overridden

Source: <http://www.coreservlets.com/hadoop-tutorial/>

# MapReduce

**Job – execution of map and reduce functions to accomplish a task**

Equal to Java's main

**Task – single Mapper or Reducer:**

Performs work on a fragment of data

**1. Configure the Job:** Specify Input, Output, Mapper, Reducer and Combiner

**2. Implement Mapper :**Input is text – e.g. a line

Tokenize the text and emit first character with a count of 1 - <token, 1>

**3. Implement Reducer** Sum up counts for each letter

Write out the result to HDFS

**4. Run the job**

# 1: Configure Job

- **Job class**

- Encapsulates information about a job
- Controls execution of the job

- **A job is packaged within a jar file**

- Hadoop Framework distributes the jar on your behalf
- Needs to know which jar file to distribute
- The easiest way to specify the jar that your job resides in is by calling `job.setJarByClass`
- Hadoop will locate the jar file that contains the provided class

## **Note: Hadoop IO Classes**

- **Hadoop uses it's own serialization mechanism for writing data in and out of network, database or files**
  - Optimized for network serialization
  - A set of basic types is provided
  - Easy to implement your own
- **org.apache.hadoop.io package**
  - LongWritable for Long
  - IntWritable for Integer
  - Text for String
  - Etc...



# 1: Configure Job - Input

```
TextInputFormat.addInputPath(job, new Path(args[0]));  
job.setInputFormatClass(TextInputFormat.class);
```

- **Can be a file, directory or a file pattern**
  - Directory is converted to a list of files as an input
- **Input is specified by implementation of InputFormat - in this case TextInputFormat**
  - Responsible for creating splits and a record reader
  - Controls input types of key-value pairs, in this case LongWritable and Text
  - – File is broken into lines, mapper will receive 1 line at a time

# 1: Configure Job – Output

```
TextOutputFormat.setOutputPath(job, new Path(args[1]));  
job.setOutputFormatClass(TextOutputFormat.class);
```

- **OutputFormat defines specification for outputting data from Map/Reduce job**
- job utilizes an implementation of OutputFormat: TextOutputFormat
  - Define output path where reducer should place its output
- If path already exists then the job will fail
  - Each reducer task writes to its own file
- By default a job is configured to run with a single reducer
  - Writes key-value pair as plain text

# 1: Configure Job – Output

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

- **Specify the output key and value types for both mapper and reducer functions**
  - – Many times the same type
  - – If types differ then use
    - setMapOutputKeyClass()
    - setMapOutputValueClass()

## 2: Implement Mapper

### **Class has 4 Java Generics\*\* parameters**

- (1) input key (2) input value (3) output key (4) output value
- Input and output utilizes hadoop's IO framework

- org.apache.hadoop.io

- **Your job is to implement map() method**

- Input key and value
- Output key and value
- Logic is up to you

- **map() method injects Context object, use to:**

- Write output
- Create your own counters

- \*\*Java Generics provide a mechanism to abstract Java types. To learn more visit: <http://docs.oracle.com/javase/tutorial/extra/generics/index.html>

### 3: Implement Reducer

- **Analogous to Mapper – generic class with four types**
  - (1) input key (2) input value (3) output key (4) output value
  - The output types of map functions must match the input types of reduce function
- **In this case Text and IntWritable**
  - Map/Reduce framework groups key-value pairs produced by mapper by key
- **For each key there is a set of one or more values**
- **Input into a reducer is sorted by key**
- **Known as Shuffle and Sort**
  - Reduce function accepts key->setOfValues and outputs key/value pairs
- **Also utilizes Context object (similar to Mapper)**

### **3: Reducer as a Combiner**

- **Combine data per Mapper task to reduce amount of data transferred to reduce phase**
- **Reducer can very often serve as a combiner**
  - Only works if reducer's output key-value pair types are the same as mapper's output types
- **Combiners are not guaranteed to run**
  - Optimization only
  - Not for critical logic

## 4: Run Count Job

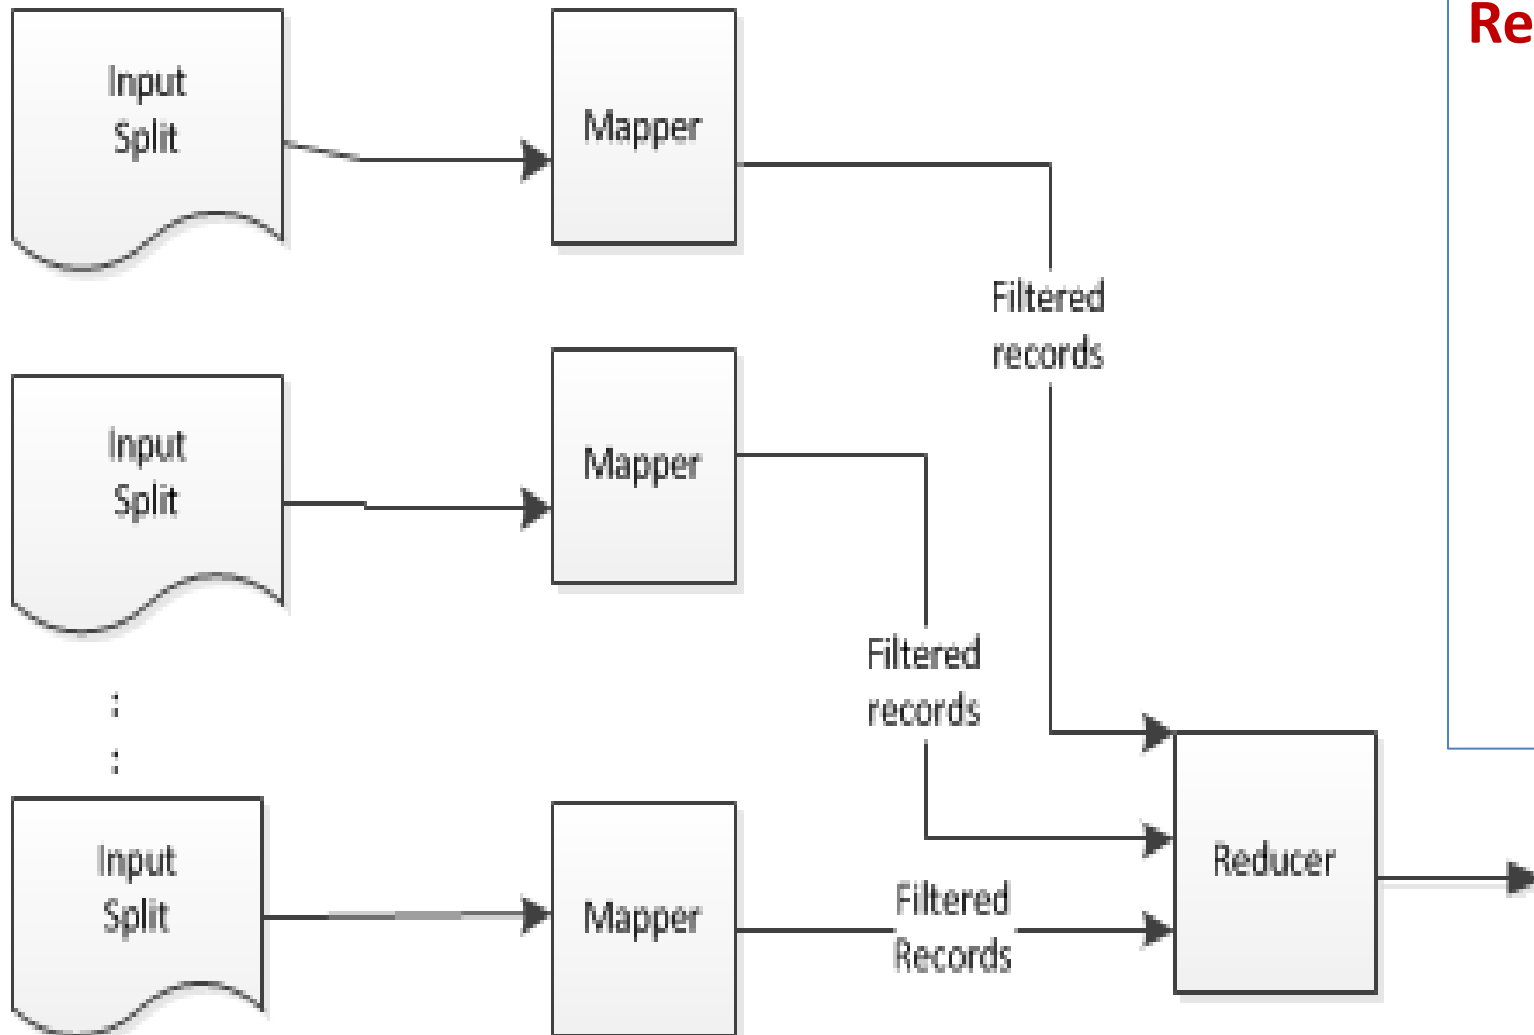
- **\*\*HANDS-ON\*\***

# Pattern: Filtering

- Most **basic** pattern.
- Use case: Filter out records that are of no interest.
- Why?
  - Recall MapReduce intermediate sorting/shuffling is I/O heavy
  - Want to reduce dataset as much as possible in the map phase.



# Filtering



**Recall Mapper signature:**

```
Map(k,v) {  
    // filter  
    if ( $f(k,v)$ )  
    {  
        :  
        emit()  
    }  
}
```

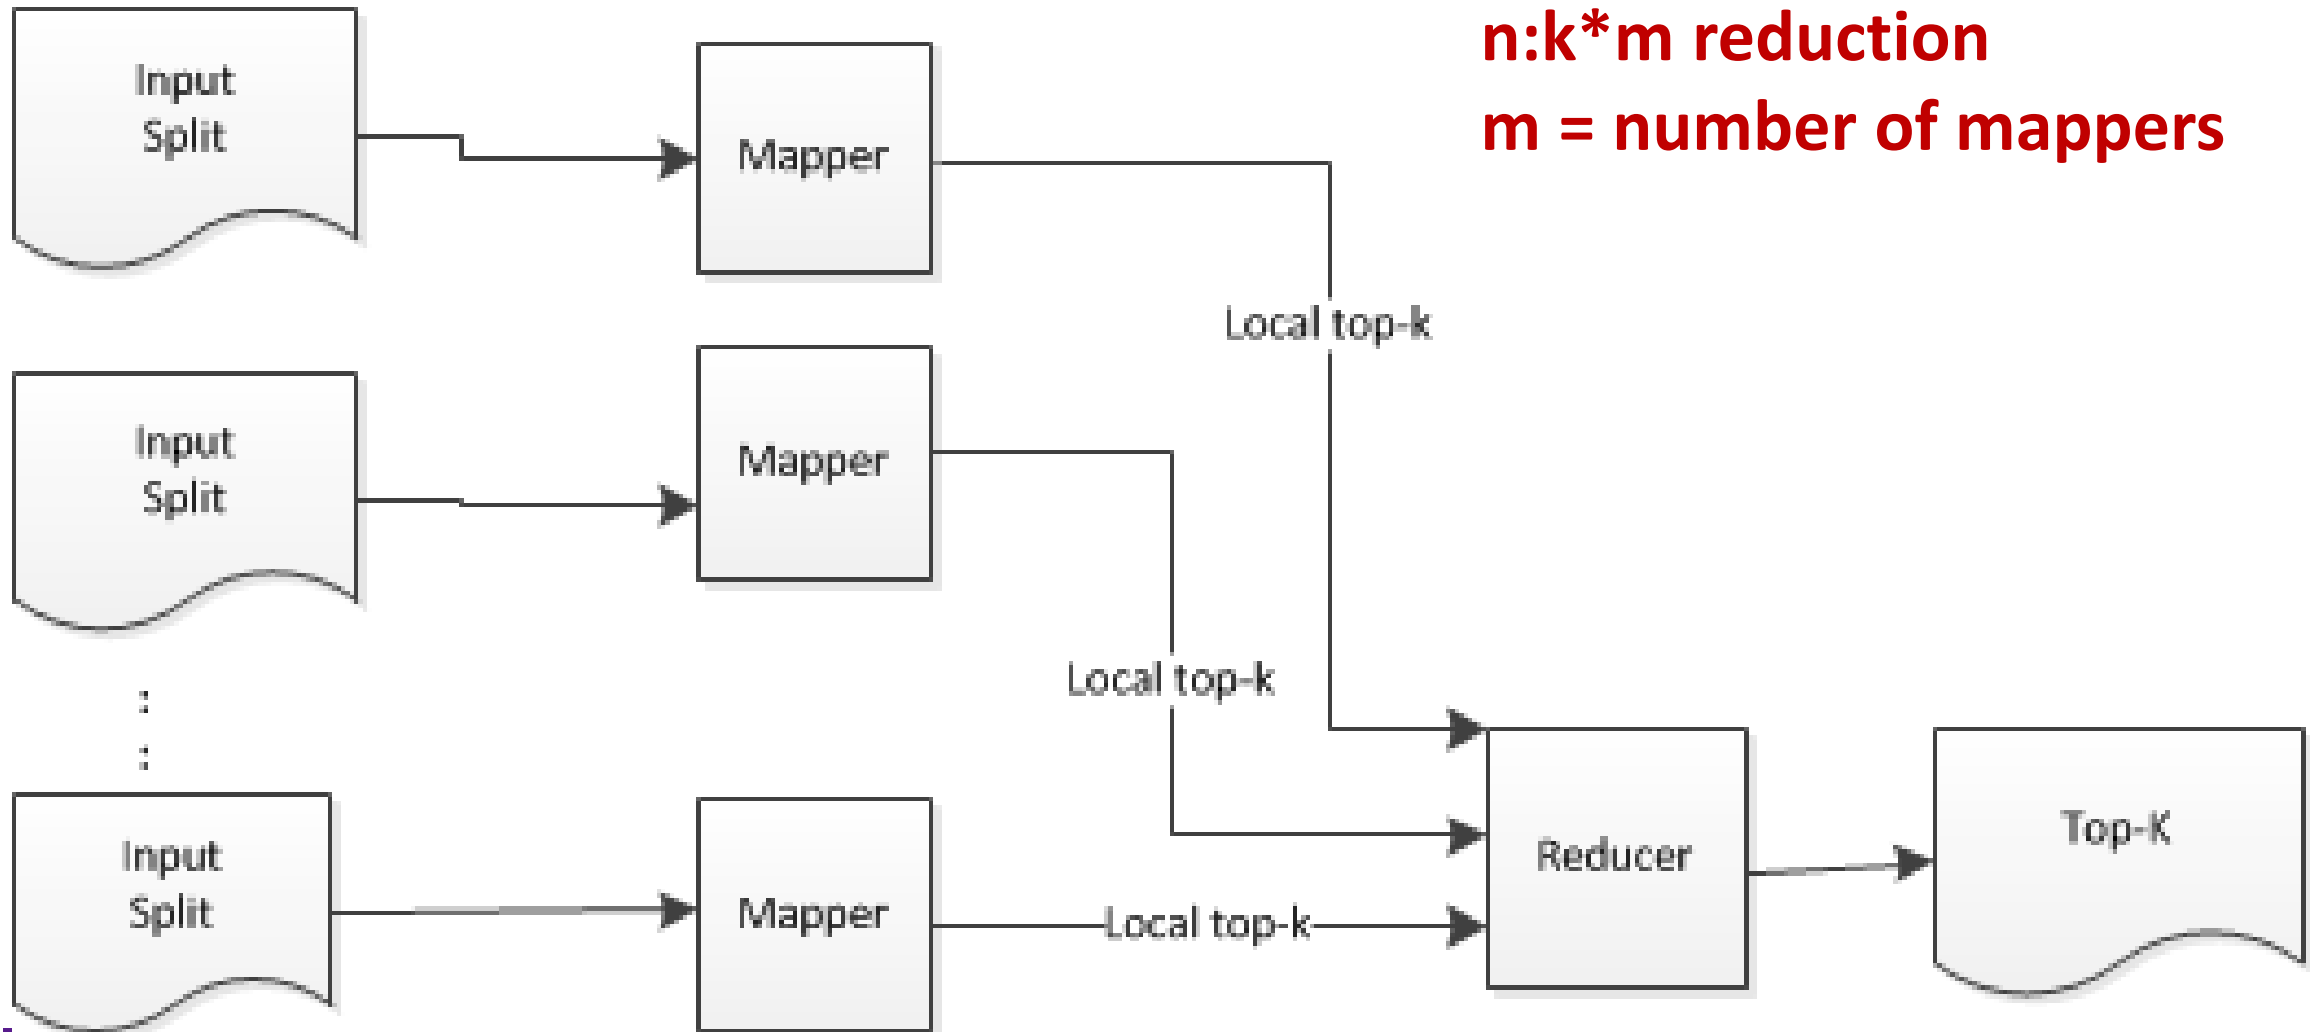
**QUESTIONS?**

# Top 10

Select the top  $k$  items from a dataset, no matter how large the data.

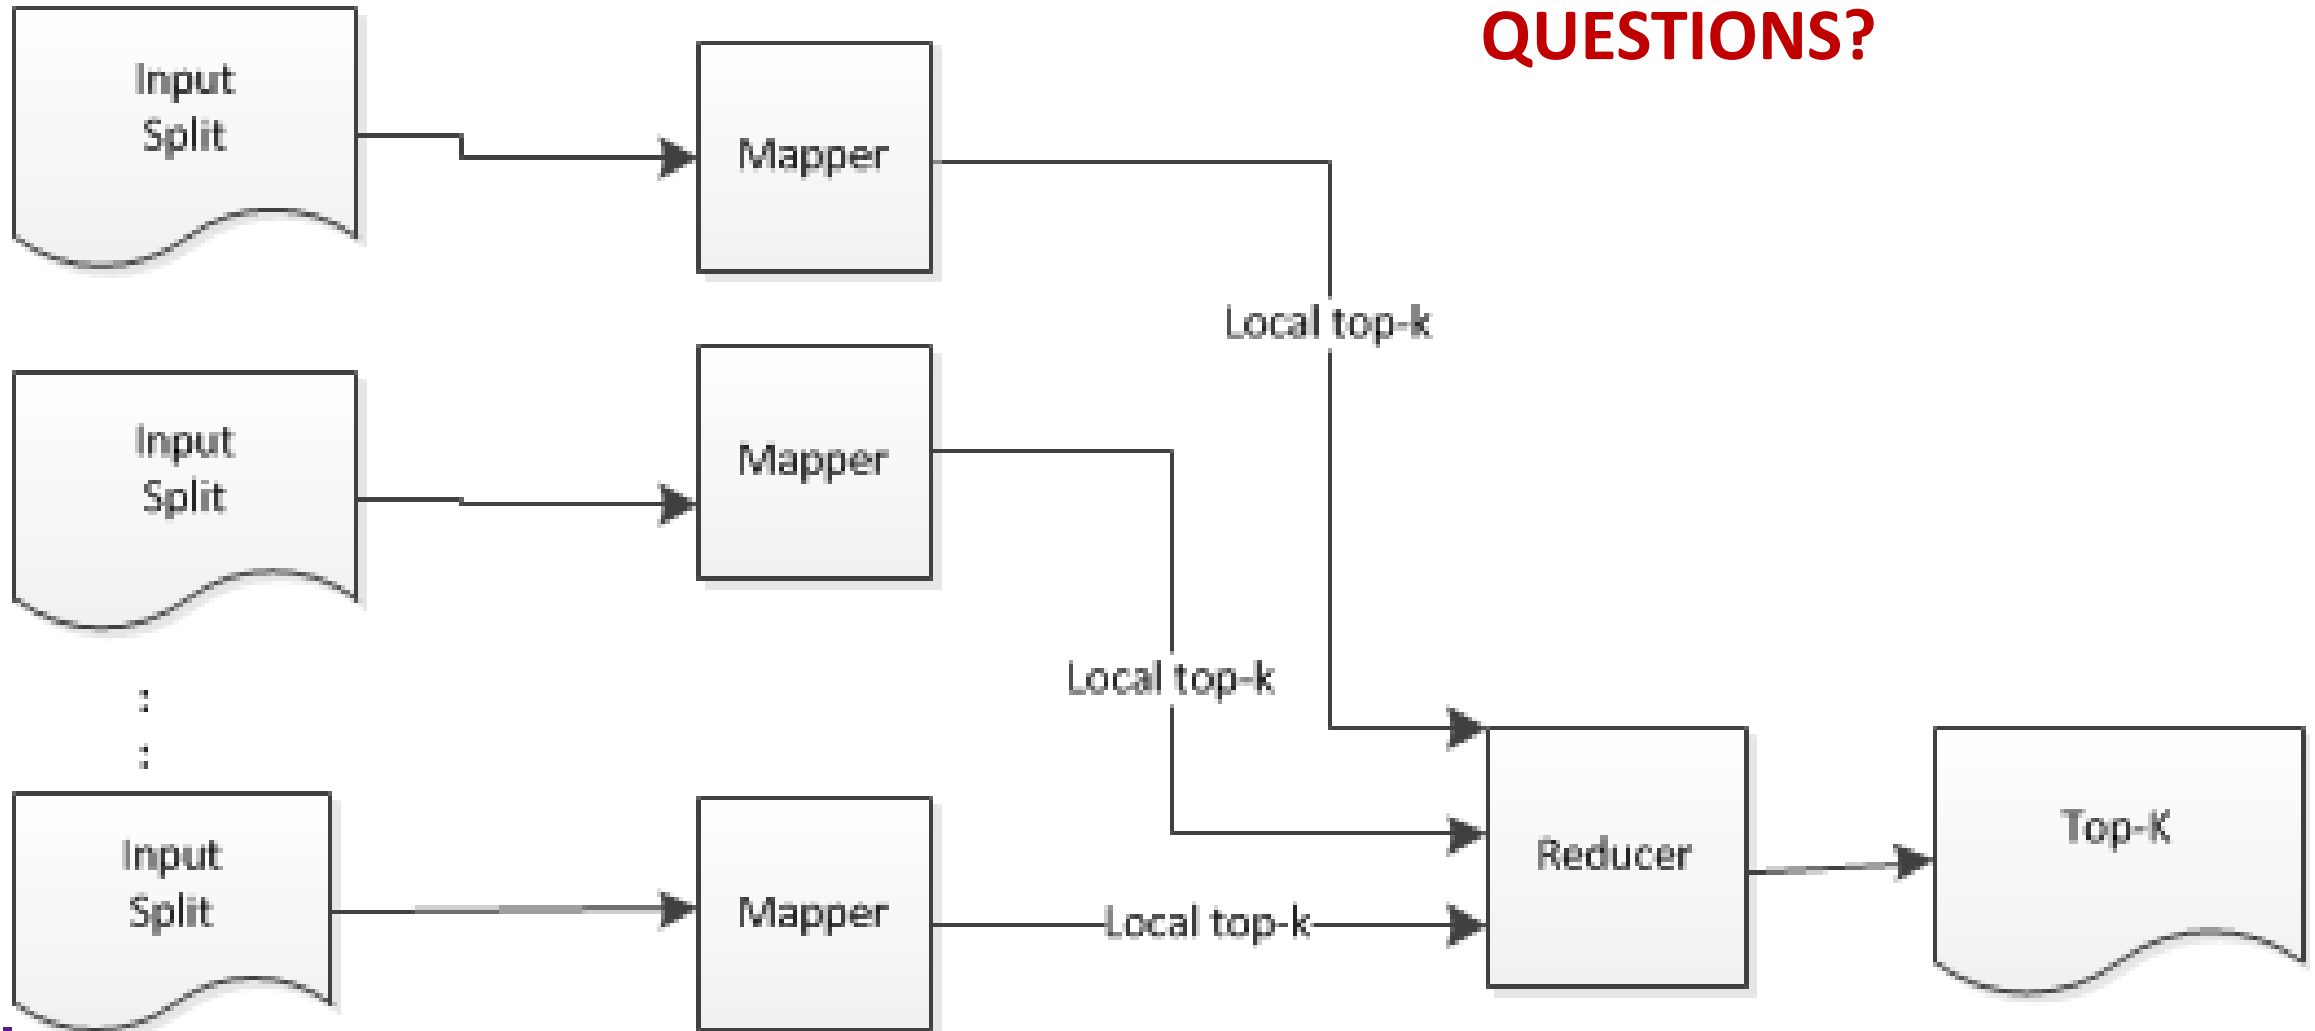
- Need to evaluate ALL the data and rank according to the top- $k$  criteria. How do we do this in MapReduce?

# Top 10



# Top 10

QUESTIONS?



# Pattern: Binning

Move records into bins/categories, irrespective of the order of records

Recall the review slide on MapReduce:

MapReduce has a partitioner class, which does the exact same job as binning.

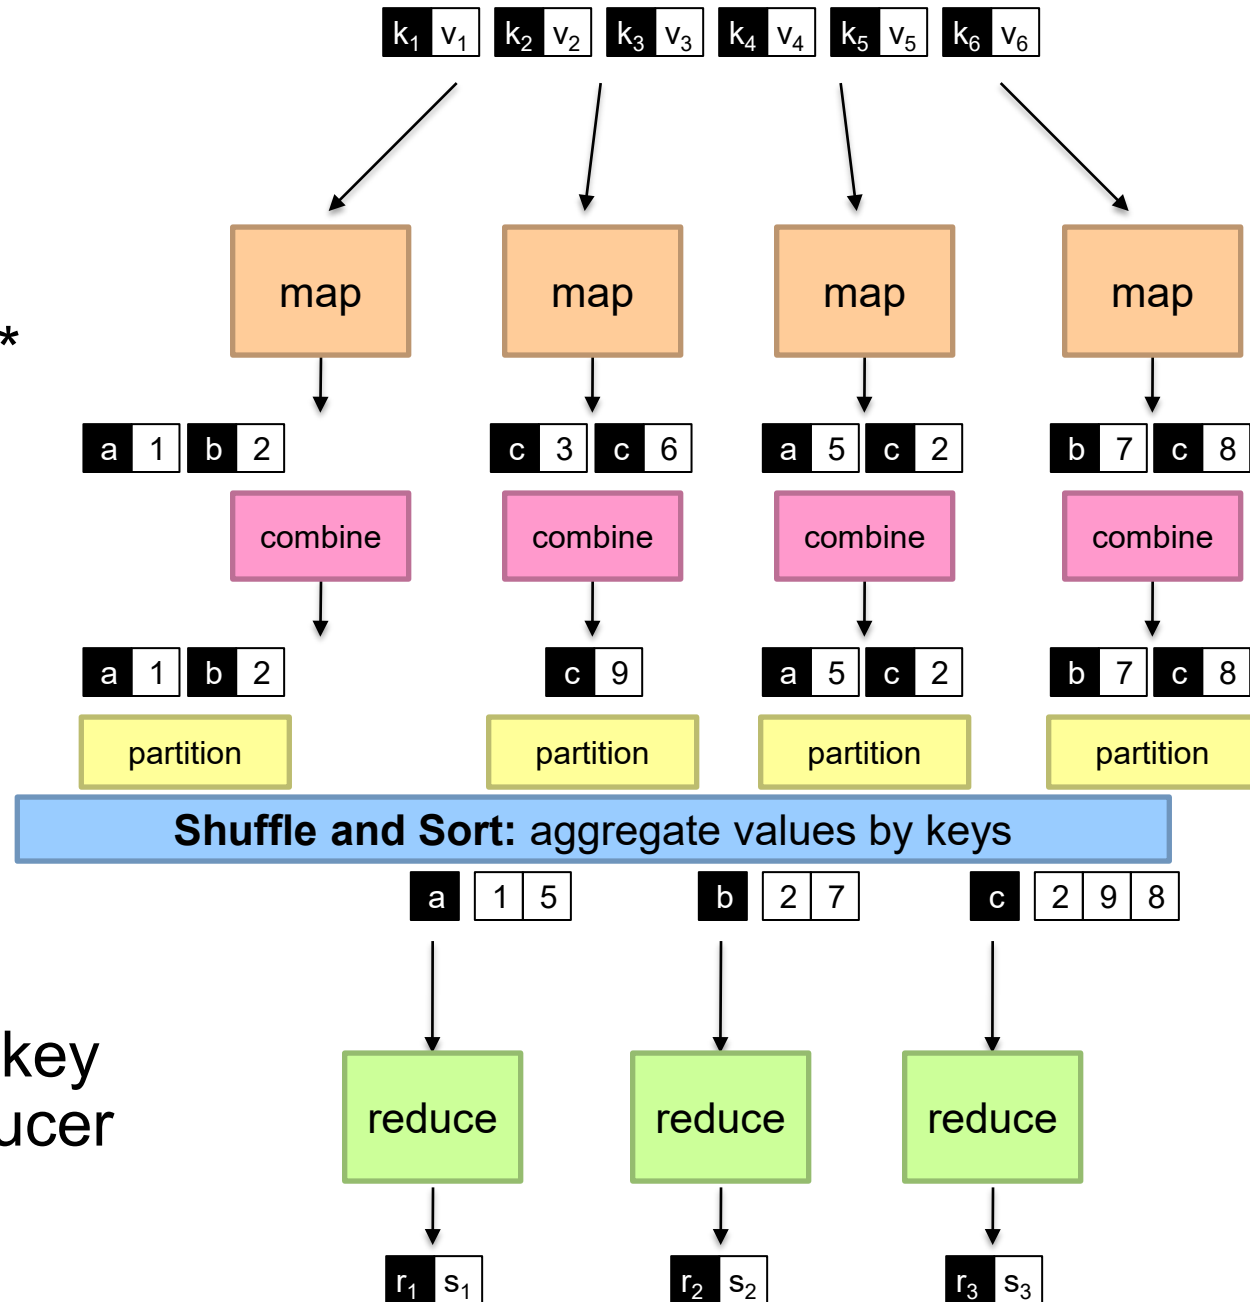
So why use binning?

**I/O**

# Binning

**map**  $(k, v) \rightarrow \langle k', v' \rangle^*$

**reduce**  $(k', v') \rightarrow \langle k', v' \rangle^*$



All values with the same key  
are sent to the same reducer

# Pattern: Binning

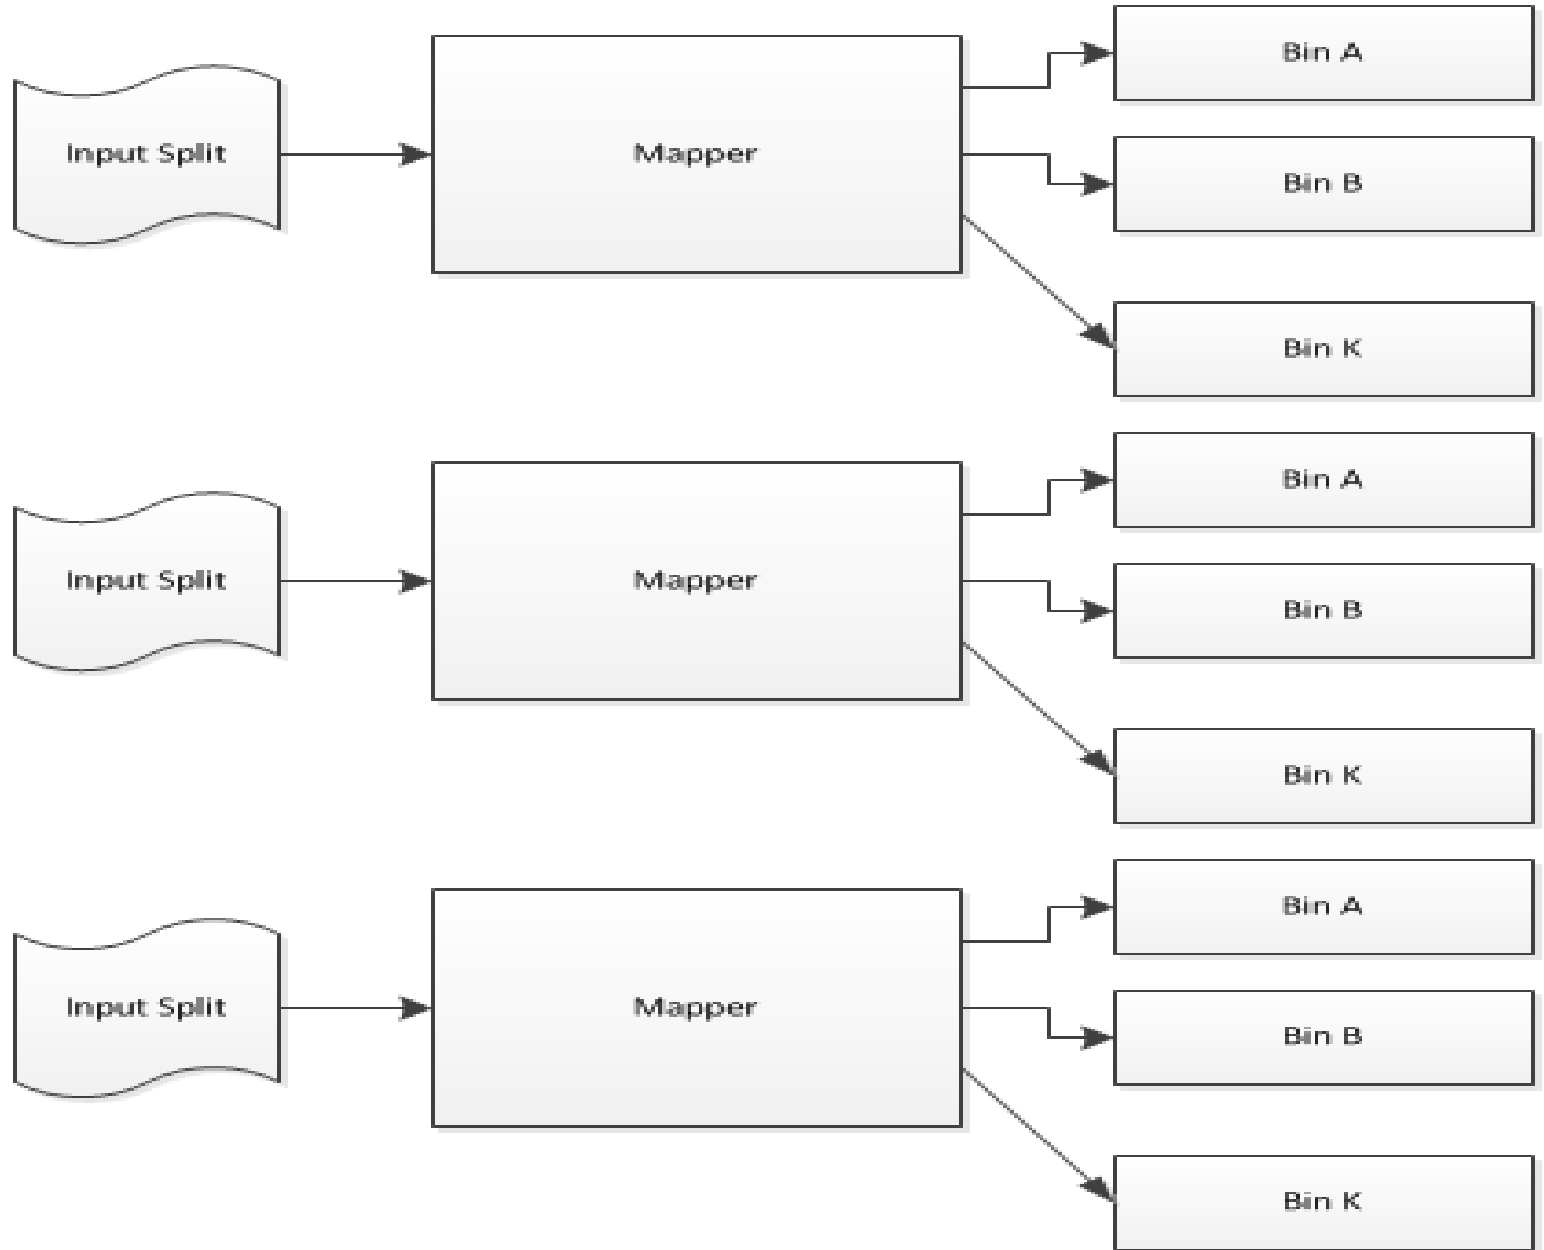
Move records into bins/categories, irrespective of the order of records

So why use binning?

I/O

# Binning


- intermediate files out of the Mapper
- No need for a reducer
- All output files can
- Just be concatenated



**QUESTIONS?**



# Pattern: BFS – Breadth First Search



BFS = general technique for traversing a graph.

BFS on a graph with  $n$  vertices and  $m$  edges:  $O(n + m)$

## Algorithm:

– Input: Simple Connected directed graph with ‘ $n$ ’ vertices and the node to be searched.

– Output: if node is found “Yes” is printed and the corresponding path is displayed  
else “No” is printed.

## Why do you care?

### Shortest Path

- BFS Guarantees to find the shortest path to the destined node if it exists in the graph.


# BFS – Breadth First Search

Graph is represented as adjacency list.

- **Key:** Node ID
- **Value:** EDGES | DISTANCE\_FROM\_SOURCE | COLOR |

Where EDGES is a comma delimited list of the ids of the nodes that are connected to this node. in the beginning, we do not know the distance and will use Integer. MAX\_VALUE for marking "unknown". Color tells us whether or not we've seen the node before, so this starts off as white.

<u>Key</u>	<u>Value</u>
1	2,3,4 0 GRAY
2	5,6 Integer.MAX_VALUE WHITE
3	NULL Integer.MAX_VALUE WHITE
4	7,8 Integer.MAX_VALUE WHITE
5	8 Integer.MAX_VALUE WHITE
:	




# BFS – Breadth First Search

## Map:

- For each gray node, the mappers emit a new gray node, with distance = distance + 1. they also then emit the input gray node, but colored red. (we're done with it.)  
Mappers also emit all non-gray nodes, with no change. so, the output of the first map iteration would be:

Key	Value
1	2,3,4 0 RED
2	NULL 1 BLUE
3	NULL 1 BLUE
4	NULL 1 BLUE
2	1,5,6 Integer.MAX_VALUE WHITE
3	NULL Integer.MAX_VALUE WHITE
4	7,8 Integer.MAX_VALUE WHITE
5	8 Integer.MAX_VALUE WHITE



# BFS – Breadth First Search

## Reduce:

Reducers, receives all data for a given key:

in this case it means that they receive the data for all "copies" of each node.

e.g, the reducer that receives the data for key = 2 gets:

2      NULL|1|BLUE|


2      1,5,6|Integer.MAX\_VALUE|WHITE|

The reducers job is to take all this data and construct a new node. Output of Reducer:

2      1,5,6|1|BLUE


Then?

Repeat MAP/REDUCE sequence until node is found, or n times, where n=# nodes




# PageRank

## Graph Algorithm



# PageRank



- Jimmy Lin and Michael Schatz.

## [Design Patterns for Efficient Graph Algorithms in MapReduce.](#)

*Proceedings of the 2010 Workshop on Mining and Learning with Graphs Workshop (MLG-2010), July 2010, Washington, D.C.*


# PageRank

## The basic algorithm:

At each iteration;  
evenly distribute a  
node's probability mass  
to its neighbors ; until  
convergence

## The hard part:

How do you do this in  
MR?



# PageRank


- **The idea**
- 1. Represent the graph as an adjacency list
- 2. Partition the graph using Hash functions (binning)
- 3. At each mapper: **emit** the node and it's entire adjacency list (need to preserve the graph structure)
- 4. At each reducer, reconstruct the node and it's neighbors, compute incoming probability mass
- 5. Repeat until converge



# PageRank

```
1: class MAPPER
2:   method MAP(id  $n$ , vertex  $N$ )
3:      $p \leftarrow N.PAGERANK / |N.ADJACENCYLIST|$ 
4:     EMTT(id  $n$ , vertex  $N$ )
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
6:       EMTT(id  $m$ , value  $p$ )

1: class REDUCER
2:   method REDUCE(id  $m$ , [ $p_1, p_2, \dots$ ])
3:      $M \leftarrow \emptyset$ 
4:     for all  $p \in [p_1, p_2, \dots]$  do
5:       if ISVERTEX( $p$ ) then
6:          $M \leftarrow p$ 
7:       else
8:          $s \leftarrow s + p$ 
9:      $M.PAGERANK \leftarrow s$ 
10:    EMTT(id  $m$ , vertex  $M$ )
```



- In the map phase we evenly divide up each vertex's PageRank mass and pass each piece along outgoing edges to neighbors.
- In the reduce phase PageRank contributions are summed up at each destination vertex.
- Each MapReduce job corresponds to one iteration of the algorithm.